

ACS Technologies Limited
Technology & Innovations Division


Gap Register ↔ AI-Adjusted Timeline Crosswalk

A pin-to-pin cross-reference of the 14 plan gaps in the Director-approved Aug–Sep 2026 closure plan against the matching task lines in the 6-June AI-Adjusted Project Timeline — checked for effort agreement, dependency agreement, and whether a corresponding task exists at all.

DOCUMENT REFERENCE	ACSTL/DoTI/DAPE/DA/2026/002
CROSS-REFERENCES	ACSTL/DoTI/DAPE/DA/2026/001 (Plan Evaluation, 07 Aug 2026)
SOURCE A	DA_gaps.pdf — Plan Evaluation Report, 07 Aug 2026
SOURCE B	AI-Adjusted Project Timeline, 06 Jun 2026
PREPARED FOR	Chandra Sekhar, Backend Lead
GAPS CROSS-CHECKED	14 of 14

01 — RESULT AT A GLANCE

Of the 14 gaps in the plan's Section 02 register, 11 map cleanly onto one or more AI-Adjusted Timeline tasks with matching scope and dependency order. Three do not — one has no timeline task at all, and two carry effort figures that read as inconsistent once the “AI effort” and “person-day” units are converted to the same basis.



02 — HOW TO READ A CARD

Each gap card sets the plan's own numbers (person-days, dependency reference) against the timeline's matching task line (AI-effort symbol, note). The verdict chip marks the outcome for that specific pairing.

- Aligned — same scope, same dependency, effort converts cleanly
- Watch — task exists but effort basis or scope needs a second look
- Mismatch — no corresponding timeline task found

03 — GAP-BY-GAP CROSSWALK

All 14 gaps, in the plan's original G1–G14 order. PD figures are the plan's own estimates; AI-effort figures and task IDs are quoted directly from the timeline document.

G13D visualisation scaffold (Cesium / terrain)		FRONTEND	ALIGNED
PLAN REGISTER — SOURCE A	TIMELINE — SOURCE B		
10 PDperson-days · dep: none Milestone target: 31 Aug, “usable scaffold.”	P3-01 + P3-02 XL (3–4 days) + L (1.5 days) ≈ 4.5–5.5 days Timeline flags Cesium as “the dominant cost — a 600k-line library with its own paradigm.” Both documents independently isolate the same task as the largest line item in their respective phase.		
Match quality: Strong. Both sources isolate Cesium integration as the long pole. 10 PD (human-paced) vs 4.5–5.5 AI-adjusted days is consistent with the timeline's own stated 3–4× bottleneck multiplier for novel third-party libraries.			

G2 Telemetry replay player (3D + timeline) FRONTEND ALIGNED	
PLAN REGISTER — SOURCE A 5 PD person-days · dep: G1	TIMELINE — SOURCE B P3-04 M (4h) · dep: P3-01
Match quality: Strong. Dependency chain agrees exactly (G2→G1 mirrors P3-04→P3-01). The effort gap (5 PD vs 4h) is the expected AI-adjusted compression for a scoped UI component, not a scope disagreement.	
G3 Drone Inventory knowledge-base workflow BOTH ALIGNED	
PLAN REGISTER — SOURCE A 8 PD person-days · dep: none	TIMELINE — SOURCE B P4-01–P4-05, P4-07 S/M cluster ≈ 1 focused day
Match quality: Strong on scope — the timeline's own note confirms “P4-01–P4-05 and P4-07 can be completed in a single focused day,” matching the plan's “Both” track split. Elasticsearch (P4-06) is correctly excluded here — it is G3's sibling gap, not G3 itself; see Section 04.	
G4 Analyst background worker (RabbitMQ) BACKEND ALIGNED	
PLAN REGISTER — SOURCE A 6 PD person-days · dep: none Milestone: 21 Aug — “Analyst job worker end-to-end.”	TIMELINE — SOURCE B P5-03 L (1.5 days) · “New infra; Docker Compose changes”
Match quality: Strong, with a naming caveat — the plan calls this RabbitMQ; the timeline implements it as Celery-with-a-broker (same infrastructure layer). Action A4 (job-state schema: PostgreSQL table vs RabbitMQ-native) should be resolved against the Celery worker-directory structure the timeline assumes.	
G5 YOLOv8 object-detection inference BACKEND WATCH	
PLAN REGISTER — SOURCE A 8 PD person-days · dep: G4 Director evaluation flags this as the single highest-uncertainty item and requires a pretrained-model baseline spike (Action A2, due 12 Aug) before the estimate is trusted.	TIMELINE — SOURCE B P5-04 L (2 days) · “ML model setup + GPU config” Timeline's own calibration table rates ML pipeline work at only 3× speedup — the lowest multiplier alongside SITL and third-party integration — and separately lists “GPU/ML environment” under “What AI Does Not Accelerate.”
Match quality: Task exists and dependency agrees (G5→G4 mirrors P5-04→P5-03). Watch: the 2-day AI-adjusted figure assumes GPU/CUDA is already provisioned — the timeline's own bottleneck table says environment setup is NOT accelerated. If Action A2's baseline spike shows mAP below 0.4, neither document's estimate holds.	

G6 Video-frame extraction pipeline (ffmpeg) BACKEND ALIGNED	
PLAN REGISTER — SOURCE A 5 PD person-days · dep: G4 Plan evaluation flags ffmpeg as absent from the original tech-stack document (Action A3).	TIMELINE — SOURCE B P5-06 M (4h) · dep: P5-03
Match quality: Strong. Both documents name it a Celery/worker-dependent task. The timeline's silence on the ffmpeg tech-stack question means Action A3 (pin version, add to Dockerfile) is still open regardless of task-level agreement.	
G7 Change-detection module BACKEND ALIGNED	
PLAN REGISTER — SOURCE A 6 PD person-days · dep: G5, G6	TIMELINE — SOURCE B P5-05 L (1.5 days) · dep: P5-03
Match quality: Scope matches ("pixel-diff algorithm" = change detection). The plan lists G7 as depending on G5+G6, but P5-05 lists only P5-03 — pixel-diff detection doesn't structurally need YOLOv8 output, so the plan's dependency is likely conservative rather than the timeline being wrong.	
G8 MinIO object storage wiring BACKEND ALIGNED	
PLAN REGISTER — SOURCE A 4 PD person-days · dep: none Milestone: 21 Aug, paired with G4.	TIMELINE — SOURCE B P5-02 S (1.5h) · "MinIO upload endpoint"
Match quality: Strong. Both pair this with the Analyst worker milestone; no dependency conflict.	
G9 Analyst frontend workspace FRONTEND ALIGNED	
PLAN REGISTER — SOURCE A 8 PD person-days · dep: G4, G5 Plan evaluation Action A7 flags an open polling-interval decision (1s vs 5s).	TIMELINE — SOURCE B P5-07 + P5-08 M (3h) + L (1.5 days) · "job submission form" + "complex canvas rendering"
Match quality: Strong once split across the two timeline tasks — the plan's single G9 line corresponds to both the submission form (P5-07) and the detection-results canvas viewer (P5-08). Action A7's polling-interval question remains unresolved in either document.	

G10 Automated test coverage $\geq 80\%$ BOTH ALIGNED	
PLAN REGISTER — SOURCE A 6 PD person-days · dep: ongoing Plan evaluation Action A8: no current coverage baseline documented.	TIMELINE — SOURCE B P6-01 L (1.5 days) · “Full pytest coverage pass”
Match quality: Strong. Note for the reader: the backend test suite has since grown substantially past this timeline's 6-June baseline (883 tests passing as of the current session) — Action A8's “run pytest-cov on Day 1” is still the correct move to get a live number.	
G11 CI pipeline (lint, type-check, test, build) BOTH ALIGNED	
PLAN REGISTER — SOURCE A 4 PD person-days · dep: none Plan explicitly schedules this mid-plan, once G1 and G4 exist as real code paths — a sequencing call the Director evaluation praised directly.	TIMELINE — SOURCE B P6-11 M (3h) · “GitHub Actions CI pipeline”
Match quality: Strong on task existence. Sequencing differs — the timeline places CI in Phase 6 (end of project) while the plan favours mid-plan. This is a scheduling choice, not a scope disagreement; the plan's ordering is the one to keep.	
G12 Accessibility pass (WCAG across 5 workspaces) FRONTEND WATCH	
PLAN REGISTER — SOURCE A 5 PD person-days · dep: none Plan evaluation Action A6: WCAG level unspecified — recommends 2.1 AA as the defence-GCS baseline.	TIMELINE — SOURCE B P6-06 M (4h) · “Needs human visual review”
Match quality: Task exists, but the effort figure doesn't convert on the same basis as most rows. The timeline's own bottleneck table lists WCAG visual review as fixed, human-gated work — yet P6-06 is coded M (4h) rather than the fixed-effort symbol used for equivalent bottlenecks (SITL, GPU/ML env) elsewhere in the same document. The plan's 5 PD figure is more consistent with the timeline's own bottleneck logic than the timeline's own task-table entry is.	

G13 Load testing (10+ drone telemetry + API) BACKEND ALIGNED	
PLAN REGISTER — SOURCE A 4 PD person-days · dep: G4–G8 stable Plan evaluation Action A10: add an explicit sub-100ms C2 latency target to the test plan, not just a survival check.	TIMELINE — SOURCE B P6-09 M (3h) · “Locust load test — 10 drones × telemetry + 3 missions”
Match quality: Strong on scope and dependency ordering (both gate this behind the full Analyst backend chain). Neither document states a latency NFR target inside the task description — Action A10 remains the deciding action regardless of which document schedules the work.	
G14 Frontend automated test suite (zero→first pass) FRONTEND MISMATCH	
PLAN REGISTER — SOURCE A 6 PD person-days · dep: none Plan milestones reference G14 twice: “31 Aug — partial” and “30 Sep — closed.”	TIMELINE — SOURCE B – no task found The 62-task register has no line item for a frontend (Vitest/RTL) test suite. P6-01 covers backend pytest coverage only; no equivalent frontend testing task exists across Phases 1–6.
Match quality: No corresponding task. This is a genuine omission in the timeline document, independent of the Director’s plan-evaluation findings — it should be added explicitly (naturally paired with P6-01, likely inside Phase 6) rather than assumed to be folded into an existing line.	

04 — FINDINGS NOT VISIBLE FROM THE GAP TABLE ALONE

Cross-reading both documents surfaces three structural points that a gap-by-gap pass alone would not show.

1

G14 has no timeline task — confirmed gap in the gap-mapping itself

Every other plan gap maps to at least one timeline task; G14 maps to none. Since the Director's evaluation already flags G10 (backend coverage) as lacking a documented baseline (Action A8), the same baseline exercise is now doubly missing on the frontend side — there is neither a plan-side Definition of Done nor a timeline-side task to anchor it to.

2

G3 and Elasticsearch (P4-06) are correctly separate, but easy to conflate

The plan's G3 (8 PD, "none" dependency) matches the timeline's P4-01–P4-05/P4-07 cluster, which the timeline explicitly says fits in one day. The timeline's P4-06 (Elasticsearch, XL 3–4 days) is a distinct, harder task that the plan's Gap Register does not list as its own line — it is implicitly bundled inside G3's "Both" track. Treating G3 as an 8 PD item risks under-budgeting the Elasticsearch portion, which the timeline's own sprint plan defers to Weeks 7–8, separate from the rest of Phase 4.

3

G12's accessibility pass is coded as acceleratable when the timeline's own bottleneck table says it isn't

This is an internal inconsistency inside the AI-Adjusted Timeline document itself, not just a mismatch against the plan: its bottleneck table lists WCAG visual review as fixed-duration and human-gated, in the same category as SITL and GPU/ML environment setup — both coded 1× (fixed) in the effort table. P6-06 is coded M (4h) instead. The plan's 5 PD figure for G12 is the more internally consistent number to carry forward.

05 — FULL CROSSWALK TABLE

Flat reference view of every row above, for copy-paste into a tracker.

GAP	DESCRIPTION	PLAN PD	PLAN DEP	TIMELINE TASK(S)	AI EFFORT	VERDICT
G1	3D visualisation scaffold	10	—	P3-01, P3-02	4.5–5.5d	Aligned
G2	Telemetry replay player	5	G1	P3-04	4h	Aligned
G3	Inventory knowledge-base workflow	8	—	P4-01–05, 07	~1d	Aligned
G4	Analyst background worker	6	—	P5-03	1.5d	Aligned
G5	YOLOv8 inference	8	G4	P5-04	2d	Watch
G6	Video-frame extraction (ffmpeg)	5	G4	P5-06	4h	Aligned
G7	Change-detection module	6	G5, G6	P5-05	1.5d	Aligned
G8	MinIO object storage wiring	4	—	P5-02	1.5h	Aligned
G9	Analyst frontend workspace	8	G4, G5	P5-07, P5-08	3h + 1.5d	Aligned
G10	Automated test coverage ≥80%	6	ongoing	P6-01	1.5d	Aligned
G11	CI pipeline	4	—	P6-11	3h	Aligned
G12	Accessibility pass (WCAG)	5	—	P6-06	4h	Watch
G13	Load testing	4	G4–G8	P6-09	3h	Aligned
G14	Frontend automated test suite	6	—	—	—	Mismatch

Prepared as a working cross-reference of ACSTL/DoTI/DAPE/DA/2026/001 (Director Plan Evaluation, 07 Aug 2026) against the AI-Adjusted Project Timeline (6 Jun 2026). Verdicts reflect textual agreement between the two source documents only — they do not independently re-estimate effort or re-verify implementation status against the current codebase.

ACS Technologies Limited
Technology & Innovations Division

Gap Register ↔ AI-Adjusted Timeline Crosswalk

A pin-to-pin cross-reference of the 14 plan gaps in the Director-approved Aug–Sep 2026 closure plan against the matching task lines in the 6-June AI-Adjusted Project Timeline — checked for effort agreement, dependency agreement, and whether a corresponding task exists at all.

DOCUMENT REFERENCE	ACSTL/DoTI/DAPE/DA/2026/002
CROSS-REFERENCES	ACSTL/DoTI/DAPE/DA/2026/001 (Plan Evaluation, 07 Aug 2026)
SOURCE A	DA_gaps.pdf — Plan Evaluation Report, 07 Aug 2026
SOURCE B	AI-Adjusted Project Timeline, 06 Jun 2026
PREPARED FOR	Chandra Sekhar, Backend Lead
GAPS CROSS-CHECKED	14 of 14

01 — RESULT AT A GLANCE

Of the 14 gaps in the plan's Section 02 register, 11 map cleanly onto one or more AI-Adjusted Timeline tasks with matching scope and dependency order. Three do not — one has no timeline task at all, and two carry effort figures that read as inconsistent once the “AI effort” and “person-day” units are converted to the same basis.

11

Gaps aligned — task exists, scope and dependency chain match

2

Gaps to watch — task exists, effort basis diverges

1

Gap with no timeline task — G14 frontend suite

02 — HOW TO READ A CARD

Each gap card sets the plan's own numbers (person-days, dependency reference) against the timeline's matching task line (AI-effort symbol, note). The verdict chip marks the outcome for that specific pairing.

- Aligned** — same scope, same dependency, effort converts cleanly
- Watch** — task exists but effort basis or scope needs a second look
- Mismatch** — no corresponding timeline task found

03 — GAP-BY-GAP CROSSWALK

All 14 gaps, in the plan's original G1–G14 order. PD figures are the plan's own estimates; AI-effort figures and task IDs are quoted directly from the timeline document.

G13D visualisation scaffold (Cesium / terrain)		FRONTEND	ALIGNED
PLAN REGISTER — SOURCE A	TIMELINE — SOURCE B		
10 PDperson-days · dep: none Milestone target: 31 Aug, “usable scaffold.”	P3-01 + P3-02 XL (3–4 days) + L (1.5 days) ≈ 4.5–5.5 days Timeline flags Cesium as “the dominant cost — a 600k-line library with its own paradigm.” Both documents independently isolate the same task as the largest line item in their respective phase.		
Match quality: Strong. Both sources isolate Cesium integration as the long pole. 10 PD (human-paced) vs 4.5–5.5 AI-adjusted days is consistent with the timeline's own stated 3–4× bottleneck multiplier for novel third-party libraries.			

G2 Telemetry replay player (3D + timeline) FRONTEND ALIGNED	
PLAN REGISTER — SOURCE A	TIMELINE — SOURCE B
5 PD person-days · dep: G1	P3-04 M (4h) · dep: P3-01
Match quality: Strong. Dependency chain agrees exactly (G2→G1 mirrors P3-04→P3-01). The effort gap (5 PD vs 4h) is the expected AI-adjusted compression for a scoped UI component, not a scope disagreement.	
G3 Drone Inventory knowledge-base workflow BOTH ALIGNED	
PLAN REGISTER — SOURCE A	TIMELINE — SOURCE B
8 PD person-days · dep: none	P4-01–P4-05, P4-07 S/M cluster ≈ 1 focused day
Match quality: Strong on scope — the timeline's own note confirms “P4-01–P4-05 and P4-07 can be completed in a single focused day,” matching the plan's “Both” track split. Elasticsearch (P4-06) is correctly excluded here — it is G3's sibling gap, not G3 itself; see Section 04.	
G4 Analyst background worker (RabbitMQ) BACKEND ALIGNED	
PLAN REGISTER — SOURCE A	TIMELINE — SOURCE B
6 PD person-days · dep: none Milestone: 21 Aug — “Analyst job worker end-to-end.”	P5-03 L (1.5 days) · “New infra; Docker Compose changes”
Match quality: Strong, with a naming caveat — the plan calls this RabbitMQ; the timeline implements it as Celery-with-a-broker (same infrastructure layer). Action A4 (job-state schema: PostgreSQL table vs RabbitMQ-native) should be resolved against the Celery worker-directory structure the timeline assumes.	
G5 YOLOv8 object-detection inference BACKEND WATCH	
PLAN REGISTER — SOURCE A	TIMELINE — SOURCE B
8 PD person-days · dep: G4 Director evaluation flags this as the single highest-uncertainty item and requires a pretrained-model baseline spike (Action A2, due 12 Aug) before the estimate is trusted.	P5-04 L (2 days) · “ML model setup + GPU config” Timeline's own calibration table rates ML pipeline work at only 3× speedup — the lowest multiplier alongside SITL and third-party integration — and separately lists “GPU/ML environment” under “What AI Does Not Accelerate.”
Match quality: Task exists and dependency agrees (G5→G4 mirrors P5-04→P5-03). Watch: the 2-day AI-adjusted figure assumes GPU/CUDA is already provisioned — the timeline's own bottleneck table says environment setup is NOT accelerated. If Action A2's baseline spike shows mAP below 0.4, neither document's estimate holds.	

G6 Video-frame extraction pipeline (ffmpeg) BACKEND ALIGNED	
PLAN REGISTER — SOURCE A 5 PD person-days · dep: G4 Plan evaluation flags ffmpeg as absent from the original tech-stack document (Action A3).	TIMELINE — SOURCE B P5-06 M (4h) · dep: P5-03
Match quality: Strong. Both documents name it a Celery/worker-dependent task. The timeline's silence on the ffmpeg tech-stack question means Action A3 (pin version, add to Dockerfile) is still open regardless of task-level agreement.	
G7 Change-detection module BACKEND ALIGNED	
PLAN REGISTER — SOURCE A 6 PD person-days · dep: G5, G6	TIMELINE — SOURCE B P5-05 L (1.5 days) · dep: P5-03
Match quality: Scope matches (“pixel-diff algorithm” = change detection). The plan lists G7 as depending on G5+G6, but P5-05 lists only P5-03 — pixel-diff detection doesn't structurally need YOLOv8 output, so the plan's dependency is likely conservative rather than the timeline being wrong.	
G8 MinIO object storage wiring BACKEND ALIGNED	
PLAN REGISTER — SOURCE A 4 PD person-days · dep: none Milestone: 21 Aug, paired with G4.	TIMELINE — SOURCE B P5-02 S (1.5h) · “MinIO upload endpoint”
Match quality: Strong. Both pair this with the Analyst worker milestone; no dependency conflict.	
G9 Analyst frontend workspace FRONTEND ALIGNED	
PLAN REGISTER — SOURCE A 8 PD person-days · dep: G4, G5 Plan evaluation Action A7 flags an open polling-interval decision (1s vs 5s).	TIMELINE — SOURCE B P5-07 + P5-08 M (3h) + L (1.5 days) · “job submission form” + “complex canvas rendering”
Match quality: Strong once split across the two timeline tasks — the plan's single G9 line corresponds to both the submission form (P5-07) and the detection-results canvas viewer (P5-08). Action A7's polling-interval question remains unresolved in either document.	

G10 Automated test coverage ≥80% BOTH ALIGNED	
PLAN REGISTER — SOURCE A 6 PD person-days · dep: ongoing Plan evaluation Action A8: no current coverage baseline documented.	TIMELINE — SOURCE B P6-01 L (1.5 days) · “Full pytest coverage pass”
Match quality: Strong. Note for the reader: the backend test suite has since grown substantially past this timeline's 6-June baseline (883 tests passing as of the current session) — Action A8's “run pytest-cov on Day 1” is still the correct move to get a live number.	
G11 CI pipeline (lint, type-check, test, build) BOTH ALIGNED	
PLAN REGISTER — SOURCE A 4 PD person-days · dep: none Plan explicitly schedules this mid-plan, once G1 and G4 exist as real code paths — a sequencing call the Director evaluation praised directly.	TIMELINE — SOURCE B P6-11 M (3h) · “GitHub Actions CI pipeline”
Match quality: Strong on task existence. Sequencing differs — the timeline places CI in Phase 6 (end of project) while the plan favours mid-plan. This is a scheduling choice, not a scope disagreement; the plan's ordering is the one to keep.	
G12 Accessibility pass (WCAG across 5 workspaces) FRONTEND WATCH	
PLAN REGISTER — SOURCE A 5 PD person-days · dep: none Plan evaluation Action A6: WCAG level unspecified — recommends 2.1 AA as the defence-GCS baseline.	TIMELINE — SOURCE B P6-06 M (4h) · “Needs human visual review”
Match quality: Task exists, but the effort figure doesn't convert on the same basis as most rows. The timeline's own bottleneck table lists WCAG visual review as fixed, human-gated work — yet P6-06 is coded M (4h) rather than the fixed-effort symbol used for equivalent bottlenecks (SITL, GPU/ML env) elsewhere in the same document. The plan's 5 PD figure is more consistent with the timeline's own bottleneck logic than the timeline's own task-table entry is.	

G13 Load testing (10+ drone telemetry + API) BACKEND ALIGNED	
PLAN REGISTER — SOURCE A 4 PD person-days · dep: G4–G8 stable Plan evaluation Action A10: add an explicit sub-100ms C2 latency target to the test plan, not just a survival check.	TIMELINE — SOURCE B P6-09 M (3h) · “Locust load test — 10 drones × telemetry + 3 missions”
Match quality: Strong on scope and dependency ordering (both gate this behind the full Analyst backend chain). Neither document states a latency NFR target inside the task description — Action A10 remains the deciding action regardless of which document schedules the work.	
G14 Frontend automated test suite (zero→first pass) FRONTEND MISMATCH	
PLAN REGISTER — SOURCE A 6 PD person-days · dep: none Plan milestones reference G14 twice: “31 Aug — partial” and “30 Sep — closed.”	TIMELINE — SOURCE B – no task found The 62-task register has no line item for a frontend (Vitest/RTL) test suite. P6-01 covers backend pytest coverage only; no equivalent frontend testing task exists across Phases 1–6.
Match quality: No corresponding task. This is a genuine omission in the timeline document, independent of the Director’s plan-evaluation findings — it should be added explicitly (naturally paired with P6-01, likely inside Phase 6) rather than assumed to be folded into an existing line.	

04 — FINDINGS NOT VISIBLE FROM THE GAP TABLE ALONE

Cross-reading both documents surfaces three structural points that a gap-by-gap pass alone would not show.

1

G14 has no timeline task — confirmed gap in the gap-mapping itself

Every other plan gap maps to at least one timeline task; G14 maps to none. Since the Director's evaluation already flags G10 (backend coverage) as lacking a documented baseline (Action A8), the same baseline exercise is now doubly missing on the frontend side — there is neither a plan-side Definition of Done nor a timeline-side task to anchor it to.

2

G3 and Elasticsearch (P4-06) are correctly separate, but easy to conflate

The plan's G3 (8 PD, "none" dependency) matches the timeline's P4-01–P4-05/P4-07 cluster, which the timeline explicitly says fits in one day. The timeline's P4-06 (Elasticsearch, XL 3–4 days) is a distinct, harder task that the plan's Gap Register does not list as its own line — it is implicitly bundled inside G3's "Both" track. Treating G3 as an 8 PD item risks under-budgeting the Elasticsearch portion, which the timeline's own sprint plan defers to Weeks 7–8, separate from the rest of Phase 4.

3

G12's accessibility pass is coded as acceleratable when the timeline's own bottleneck table says it isn't

This is an internal inconsistency inside the AI-Adjusted Timeline document itself, not just a mismatch against the plan: its bottleneck table lists WCAG visual review as fixed-duration and human-gated, in the same category as SITL and GPU/ML environment setup — both coded 1× (fixed) in the effort table. P6-06 is coded M (4h) instead. The plan's 5 PD figure for G12 is the more internally consistent number to carry forward.

05 — FULL CROSSWALK TABLE

Flat reference view of every row above, for copy-paste into a tracker.

GAP	DESCRIPTION	PLAN PD	PLAN DEP	TIMELINE TASK(S)	AI EFFORT	VERDICT
G1	3D visualisation scaffold	10	—	P3-01, P3-02	4.5–5.5d	Aligned
G2	Telemetry replay player	5	G1	P3-04	4h	Aligned
G3	Inventory knowledge-base workflow	8	—	P4-01–05, 07	~1d	Aligned
G4	Analyst background worker	6	—	P5-03	1.5d	Aligned
G5	YOLOv8 inference	8	G4	P5-04	2d	Watch
G6	Video-frame extraction (ffmpeg)	5	G4	P5-06	4h	Aligned
G7	Change-detection module	6	G5, G6	P5-05	1.5d	Aligned
G8	MinIO object storage wiring	4	—	P5-02	1.5h	Aligned
G9	Analyst frontend workspace	8	G4, G5	P5-07, P5-08	3h + 1.5d	Aligned
G10	Automated test coverage ≥80%	6	ongoing	P6-01	1.5d	Aligned
G11	CI pipeline	4	—	P6-11	3h	Aligned
G12	Accessibility pass (WCAG)	5	—	P6-06	4h	Watch
G13	Load testing	4	G4–G8	P6-09	3h	Aligned
G14	Frontend automated test suite	6	—	—	—	Mismatch

Prepared as a working cross-reference of ACSTL/DoTI/DAPE/DA/2026/001 (Director Plan Evaluation, 07 Aug 2026) against the AI-Adjusted Project Timeline (6 Jun 2026). Verdicts reflect textual agreement between the two source documents only — they do not independently re-estimate effort or re-verify implementation status against the current codebase.